

Laboratorio Reglas Asociación

Pollo

20/1/2017

Reglas de Asociación

Cargamos la BD en nuestra zona de trabajo y consultamos sus dimensiones. Vemos las 2 primeras filas para ver los atributos y sus tipos

```
library(arules)
data()
data("AdultUCI")
```

```
dim(AdultUCI)
```

```
## [1] 48842    15
```

```
AdultUCI[1:2,]
```

```
##   age      workclass fnlwgt education education-num   marital-status
## 1  39      State-gov  77516 Bachelors             13   Never-married
## 2  50 Self-emp-not-inc 83311 Bachelors             13 Married-civ-spouse
##      occupation relationship race sex capital-gain capital-loss
## 1   Adm-clerical Not-in-family White Male          2174          0
## 2 Exec-managerial      Husband White Male           0          0
##   hours-per-week native-country income
## 1              40   United-States small
## 2              13   United-States small
```

De los 6 atributos continuos: 2 los eliminamos porque aportan información redundante: fnlwgt y education-num y los 4 restantes los dividimos en intervalos. Finalmente convertimos el data.frame en un conjunto de transacciones con la función as

```
AdultUCI[["fnlwgt"]] = NULL
AdultUCI[["education-num"]] = NULL
AdultUCI[["age"]] = ordered( cut ( AdultUCI[["age"]], c(0,25,45,65,100) ) ,
  labels = c ("Young", "Middle-aged", "Senior", "Old"))
AdultUCI[["hours-per-week"]] = ordered( cut ( AdultUCI[["hours-per-week"]],
  c(0,25,40,60,168) ) ,
  labels = c("Part-time", "Full-time", "Over-time", "Workaholic"))
AdultUCI[["capital-gain"]] = ordered( cut ( AdultUCI[["capital-gain"]],
  c(-Inf,0,median(AdultUCI[["capital-gain"]][AdultUCI[["capital-gain"]]>0]), Inf) ) ,
  labels = c("None", "Low", "High"))
AdultUCI[["capital-loss"]] = ordered( cut ( AdultUCI[["capital-loss"]],
  c(-Inf,0, median(AdultUCI[["capital-loss"]][AdultUCI[["capital-loss"]]>0]), Inf) ) ,
  labels = c("None", "Low", "High"))
Adult <- as(AdultUCI,"transactions")
Adult
```

```
## transactions in sparse format with
## 48842 transactions (rows) and
## 115 items (columns)
```

Vemos el resumen de la BD

```
summary(Adult)
```

```
## transactions as itemMatrix in sparse format with
## 48842 rows (elements/itemsets/transactions) and
## 115 columns (items) and a density of 0.1089939
##
## most frequent items:
##          capital-loss=None          capital-gain=None
##                46560                44807
## native-country=United-States          race=White
##                43832                41762
##                workclass=Private          (Other)
##                33906                401333
##
## element (itemset/transaction) length distribution:
## sizes
##      9      10      11      12      13
##    19    971   2067  15623  30162
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      9.00   12.00   13.00   12.53   13.00   13.00
##
## includes extended item information - examples:
##      labels variables      levels
## 1      age=Young      age      Young
## 2 age=Middle-aged      age Middle-aged
## 3      age=Senior      age      Senior
##
## includes extended transaction information - examples:
##      transactionID
## 1              1
## 2              2
## 3              3
```

Representar gráficamente la distribución de los items en las transacciones. Como en Adult cada transacción tienen un valor cada atributo/variable, usamos para probarlo la BD Epub (15729 transacciones y 936 items)


```
data(Epub)
```

```
summary(Epub)
```

```
## transactions as itemMatrix in sparse format with
## 15729 rows (elements/itemsets/transactions) and
## 936 columns (items) and a density of 0.001758755
##
## most frequent items:
## doc_11d doc_813 doc_4c6 doc_955 doc_698 (Other)
```

```
##      356      329      288      282      245      24393
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9     10     11     12
## 11615 2189  854  409  198  121  93   50   42   34   26   12
##      13     14     15     16     17     18     19     20     21     22     23     24
##      10     10      6      8      6      5      8      2      2      3      2      3
##      25     26     27     28     30     34     36     38     41     43     52     58
##       4      5      1      1      1      2      1      2      1      1      1      1
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    1.000  1.000   1.000   1.646   2.000  58.000
##
## includes extended item information - examples:
##      labels
## 1 doc_11d
## 2 doc_13d
## 3 doc_14c
##
## includes extended transaction information - examples:
##      transactionID      TimeStamp
## 10792 session_4795 2003-01-02 02:59:00
## 10793 session_4797 2003-01-02 13:46:01
## 10794 session_479a 2003-01-02 16:50:38
```


```
image(Epub)
```



Para ver gráficamente que items son los más importantes: donde el mínimo soporte será 0.1 y reducimos el

tamaño de los títulos

```
itemFrequencyPlot(Adult, support = 0.1, cex.names=0.8)
```



Usamos apriori para extraer los itemsets frecuentes con minsup 0.1. Orenamos por el valor de soporte. Inspeccionamos los 10 primeros

```
iAdult <- apriori(Adult, parameter = list(support = 0.1, target="frequent"))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          NA      0.1    1 none FALSE                TRUE         5     0.1    1
## maxlen          target  ext
##      10 frequent itemsets FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 4884
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[115 item(s), 48842 transaction(s)] done [0.03s].
## sorting and recoding items ... [31 item(s)] done [0.01s].
## creating transaction tree ... done [0.03s].
## checking subsets of size 1 2 3 4 5 6 7 8 9 done [0.10s].
## writing ... [2616 set(s)] done [0.00s].
## creating S4 object ... done [0.02s].
```

```
iAdult <- sort(iAdult, by="support")
inspect(head(iAdult, n=10))
```

```
##      items                                support
## [1] {capital-loss=None}                    0.9532779
## [2] {capital-gain=None}                    0.9173867
## [3] {native-country=United-States}          0.8974243
## [4] {capital-gain=None,capital-loss=None}    0.8706646
## [5] {race=White}                            0.8550428
## [6] {capital-loss=None,native-country=United-States} 0.8548380
## [7] {capital-gain=None,native-country=United-States} 0.8219565
## [8] {race=White,capital-loss=None}            0.8136849
## [9] {race=White,native-country=United-States}    0.7881127
## [10] {race=White,capital-gain=None}             0.7817862
```

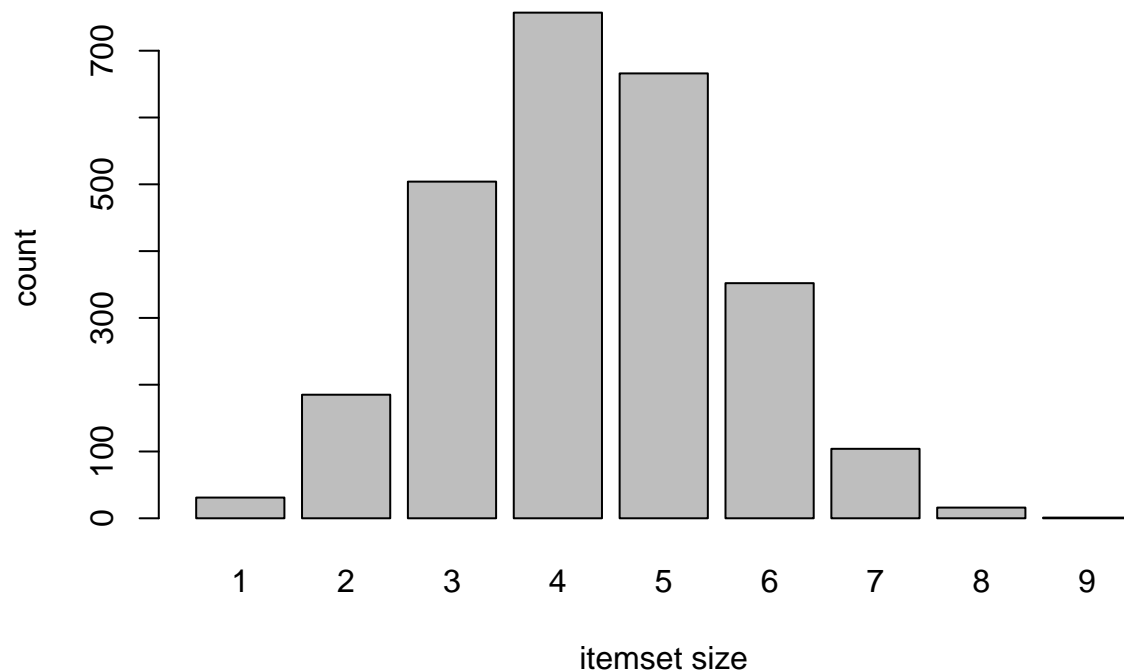
Podemos consultar el tamaño de los itemsets frecuentes (solo los 200 primeros)

```
size(iAdult)[1:200]
```

```
##      [1] 1 1 1 2 1 2 2 2 2 2 3 3 3 3 1 4 1 2 2 2 2 3 2 2 2 3 2 1 3 3 3 3 2 3 3
##      [36] 2 3 3 4 3 3 4 3 4 2 4 4 1 1 4 4 3 2 4 2 2 3 2 5 3 2 3 1 4 5 2 2 2 3 3
##      [71] 3 3 3 3 2 3 4 2 2 2 4 4 3 3 3 2 2 2 3 4 3 3 1 2 2 3 3 2 3 3 4 4 3 3 4
##     [106] 3 5 3 4 3 3 4 4 3 2 3 3 4 4 3 2 4 4 3 3 4 4 4 2 2 3 3 3 4 4 2 2 3 3 4
##     [141] 3 3 3 5 4 2 3 3 3 4 4 3 3 2 3 3 5 4 5 4 3 3 4 4 4 5 4 3 4 4 5 3 4
##     [176] 4 4 5 4 3 5 4 4 4 4 1 5 3 1 5 3 4 4 5 3 4 4 4 1 3
```

Representamos con un diagrama de barras

```
barplot(table(size(iAdult)), xlab="itemset size", ylab="count")
```



Inspeccionamos los itemsets frecuentes de tamaño 1

```
inspect(iAdult[size(iAdult)==1])
```

```
##      items                                support
## [1] {capital-loss=None}                    0.9532779
## [2] {capital-gain=None}                    0.9173867
## [3] {native-country=United-States}         0.8974243
## [4] {race=White}                           0.8550428
## [5] {workclass=Private}                      0.6941976
## [6] {sex=Male}                               0.6684820
## [7] {hours-per-week=Full-time}              0.5850907
## [8] {income=small}                           0.5061218
## [9] {age=Middle-aged}                       0.5051185
## [10] {marital-status=Married-civ-spouse}    0.4581917
## [11] {relationship=Husband}                   0.4036690
## [12] {sex=Female}                           0.3315180
## [13] {marital-status=Never-married}          0.3299824
## [14] {education=HS-grad}                     0.3231645
## [15] {age=Senior}                             0.2608616
## [16] {hours-per-week=Over-time}               0.2595307
## [17] {relationship=Not-in-family}            0.2576266
## [18] {education=Some-college}                 0.2227182
## [19] {age=Young}                              0.1971050
## [20] {education=Bachelors}                   0.1643053
## [21] {income=large}                           0.1605381
## [22] {relationship=Own-child}                  0.1552148
## [23] {marital-status=Divorced}                0.1358052
## [24] {occupation=Prof-specialty}               0.1263667
## [25] {occupation=Craft-repair}                0.1251382
## [26] {occupation=Exec-managerial}              0.1246059
## [27] {hours-per-week=Part-time}               0.1210638
## [28] {occupation=Adm-clerical}                  0.1148806
## [29] {occupation=Sales}                         0.1126899
## [30] {relationship=Unmarried}                 0.1049302
## [31] {occupation=Other-service}               0.1007944
```

Sacamos un vector lógico indicando que itemsets es maximal y mostramos los 6 primeros ordenados por su valor de soporte

```
imaxAdult <- iAdult[is.maximal(iAdult)]
inspect(head(sort(imaxAdult, by="support")))
```

```
##      items                                support
## [1] {workclass=Private,
##      race=White,
##      sex=Male,
##      capital-gain=None,
##      capital-loss=None,
##      hours-per-week=Full-time,
##      native-country=United-States}    0.1774293
## [2] {workclass=Private,
##      race=White,
##      capital-gain=None,
```

```
##      capital-loss=None,
##      hours-per-week=Full-time,
##      native-country=United-States,
##      income=small}          0.1578150
## [3] {workclass=Private,
##      race=White,
##      sex=Male,
##      capital-gain=None,
##      capital-loss=None,
##      native-country=United-States,
##      income=small}          0.1560952
## [4] {age=Middle-aged,
##      workclass=Private,
##      race=White,
##      capital-gain=None,
##      capital-loss=None,
##      hours-per-week=Full-time,
##      native-country=United-States} 0.1456124
## [5] {marital-status=Married-civ-spouse,
##      relationship=Husband,
##      race=White,
##      sex=Male,
##      capital-gain=None,
##      capital-loss=None,
##      hours-per-week=Full-time,
##      native-country=United-States} 0.1429712
## [6] {race=White,
##      sex=Female,
##      capital-gain=None,
##      capital-loss=None,
##      native-country=United-States,
##      income=small}          0.1339216
```


Sacamos un vector lógico indicando que itemsets es cerrado y mostramos los 6 primeros ordenados por su valor de soporte

```
icloAdult <- iAdult[is.closed(iAdult)]
inspect(head(sort(icloAdult, by="support")))
```

```
##      items                                support
## [1] {capital-loss=None}                    0.9532779
## [2] {capital-gain=None}                    0.9173867
## [3] {native-country=United-States}         0.8974243
## [4] {capital-gain=None,capital-loss=None}  0.8706646
## [5] {race=White}                          0.8550428
## [6] {capital-loss=None,native-country=United-States} 0.8548380
```

Podemos pintar un gráfico de barras para ver la cantidad de itemsets frecuentes, cerrados y maximales que se han generado

```
barplot( c(frequent=length(iAdult), closed=length(icloAdult), maximal=length(imaxAdult)), ylab="count",
```



Usamos apriori para extraer las reglas con mínimo soporte 0.1 y confianza 0.8 con una longitud minima de 2. Obtenemos información resumida del conjunto

```
rules <- apriori(Adult, parameter = list(support = 0.1, confidence = 0.8, minlen = 2))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.8   0.1   1 none FALSE              TRUE     5    0.1    2
## maxlen target  ext
##      10  rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 4884
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[115 item(s), 48842 transaction(s)] done [0.03s].
## sorting and recoding items ... [31 item(s)] done [0.01s].
## creating transaction tree ... done [0.02s].
## checking subsets of size 1 2 3 4 5 6 7 8 9 done [0.09s].
## writing ... [6133 rule(s)] done [0.00s].
## creating S4 object ... done [0.01s].
```

```
summary(rules)
```

```
## set of 6133 rules
##
```



```
## rule length distribution (lhs + rhs):sizes
##      2      3      4      5      6      7      8      9
## 121  637 1510 1903 1345  511   99    7
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    2.000   4.000   5.000   4.926   6.000   9.000
##
## summary of quality measures:
##      support      confidence      lift
## Min.      :0.1000   Min.      :0.8004   Min.      :0.9169
## 1st Qu.:0.1158   1st Qu.:0.8895   1st Qu.:0.9911
## Median :0.1353   Median :0.9241   Median :1.0197
## Mean    :0.1700   Mean    :0.9236   Mean     :1.2044
## 3rd Qu.:0.1890   3rd Qu.:0.9587   3rd Qu.:1.0783
## Max.    :0.8707   Max.    :1.0000   Max.     :2.9421
##
## mining info:
## data ntransactions support confidence
## Adult          48842      0.1         0.8
```

Podemos ver las reglas (lhs es el antecedente y rhs el consecuente de la regla) y sus valores para las medidas soporte, confianza y lift. También podemos ver solo los valores de las medidas de calidad

```
inspect(head(rules))
```

```
##      lhs                                rhs      support confidence      lift
## [1] {relationship=Unmarried} => {capital-loss=None} 0.1019819 0.9719024 1.0195373
## [2] {occupation=Sales}      => {race=White}      0.1005282 0.8920785 1.0433144
## [3] {occupation=Sales}      => {native-country=United-States} 0.1039679 0.9226017 1.0280552
## [4] {occupation=Sales}      => {capital-gain=None} 0.1030670 0.9146076 0.9969706
## [5] {occupation=Sales}      => {capital-loss=None} 0.1068343 0.9480378 0.9945030
## [6] {occupation=Adm-clerical} => {native-country=United-States} 0.1052373 0.9160577 1.0207632
```

```
quality(head(rules))
```

```
##      support confidence      lift
## 1 0.1019819 0.9719024 1.0195373
## 2 0.1005282 0.8920785 1.0433144
## 3 0.1039679 0.9226017 1.0280552
## 4 0.1030670 0.9146076 0.9969706
## 5 0.1068343 0.9480378 0.9945030
## 6 0.1052373 0.9160577 1.0207632
```

Podemos ordenar las reglas por el campo que más nos interese

```
rulesSorted = sort(rules, by="confidence")
inspect(head(rulesSorted))
```

```
##      lhs                                rhs      support confidence      lift
## [1] {relationship=Husband,
##      income=large}                        => {sex=Male} 0.1211662      1 1.495926
```

```
## [2] {relationship=Husband,
##      hours-per-week=Over-time}      => {sex=Male} 0.1472298      1 1.495926
## [3] {age=Senior,
##      relationship=Husband}          => {sex=Male} 0.1479874      1 1.495926
## [4] {marital-status=Married-civ-spouse,
##      relationship=Husband,
##      income=large}                  => {sex=Male} 0.1210843      1 1.495926
## [5] {relationship=Husband,
##      race=White,
##      income=large}                  => {sex=Male} 0.1111339      1 1.495926
## [6] {relationship=Husband,
##      native-country=United-States,
##      income=large}                  => {sex=Male} 0.1110724      1 1.495926
```

Seleccionar un subconjunto de reglas que cumplan una condición. Por ejemplo, seleccionamos las reglas que tenga lift > 1.2 y que en el consecuente de la regla tengan el itemset race=White

```
rulesRaceWhite <- subset(rules, subset = lhs %in% "race=White" & lift > 1.2)
inspect(head(rulesRaceWhite))
```

```
##      lhs                                rhs                                support confidence    lift
## [1] {occupation=Craft-repair,
##      race=White}                        => {sex=Male}                        0.1076532  0.9553052 1.429066
## [2] {relationship=Own-child,
##      race=White}                        => {marital-status=Never-married}    0.1168257  0.8959020 2.714999
## [3] {race=White,
##      income=large}                      => {marital-status=Married-civ-spouse} 0.1249949  0.8578053 1.872154
## [4] {race=White,
##      income=large}                      => {sex=Male}                        0.1246673  0.8555571 1.279851
## [5] {age=Young,
##      race=White}                        => {marital-status=Never-married}    0.1440154  0.8512647 2.579728
## [6] {race=White,
##      hours-per-week=Over-time} => {sex=Male}                        0.1956103  0.8239759 1.232607
```

Eliminar las reglas redundantes

```
subsetMatrix <- is.subset(rulesSorted, rulesSorted)
subsetMatrix[lower.tri(subsetMatrix, diag=T)] <- NA
redundant <- colSums(subsetMatrix, na.rm=T) >= 1
rulesPruned <- rulesSorted[!redundant] # remove redundant rules
inspect(head(rulesPruned))
```

```
##      lhs                                rhs                                support confidence    lift
## [1] {relationship=Husband,
##      income=large}                        => {sex=Male}                        0.1211662  1.0000000 1.495926
## [2] {relationship=Husband,
##      hours-per-week=Over-time} => {sex=Male}                        0.1472298  1.0000000 1.495926
## [3] {age=Senior,
##      relationship=Husband}                => {sex=Male}                        0.1479874  1.0000000 1.495926
## [4] {relationship=Husband}                => {sex=Male}                        0.4036485  0.9999493 1.495851
## [5] {age=Senior,
##      relationship=Husband}                => {marital-status=Married-civ-spouse} 0.1479669  0.9998616 2.182191
```

```
## [6] {relationship=Husband,
##      race=White,
##      hours-per-week=Full-time} => {marital-status=Married-civ-spouse} 0.1886491 0.9996745 2.181782
```

También podemos calcular para itemsets o para reglas otras medidas de calidad. Podemos calcular estas medidas para nuestras reglas podadas y añadirlas a la sección quality para que los valores de las medidas nuevas salgan también cuando inspeccionamos las reglas:


```
mInteres <- interestMeasure(rulesPruned, measure=c("hyperConfidence", "leverage", "phi", "gini"), trans
quality(rulesPruned) <- cbind(quality(rulesPruned), mInteres)
inspect(head(sort(rulesPruned, by="phi")))
```

| | lhs | rhs | support | confidence | lift |
|--------|--|--|-----------|------------|----------|
| ## [1] | {relationship=Husband} | => {marital-status=Married-civ-spouse} | 0.4034233 | 0.9993914 | 2.181164 |
| ## [2] | {relationship=Husband, race=White} | => {marital-status=Married-civ-spouse} | 0.3654232 | 0.9994400 | 2.181270 |
| ## [3] | {relationship=Husband} | => {sex=Male} | 0.4036485 | 0.9999493 | 1.495851 |
| ## [4] | {relationship=Husband, hours-per-week=Full-time} | => {marital-status=Married-civ-spouse} | 0.2147742 | 0.9995236 | 2.181453 |
| ## [5] | {age=Young} | => {marital-status=Never-married} | 0.1684820 | 0.8547834 | 2.590391 |
| ## [6] | {relationship=Husband, race=White, hours-per-week=Full-time} | => {marital-status=Married-civ-spouse} | 0.1886491 | 0.9996745 | 2.181782 |

```
library(arulesViz)
```

Utilizar la función plot para representar las reglas en función de las medidas de calidad

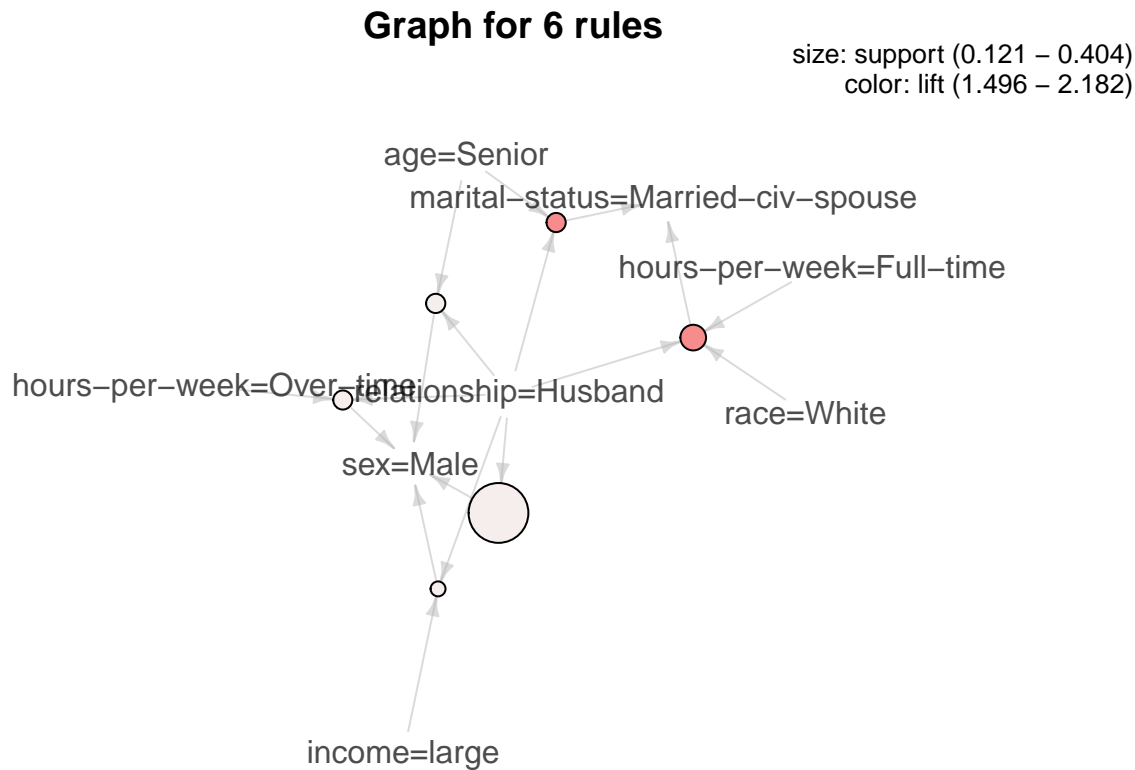
```
plot(rulesPruned)
```



Podemos modificar el tipo de gráfico generado cambiando el parámetro método de la función plot. Además, se puede modificar el gráfico cambiando los parámetros del tipo de gráfico

```
??plot # consultar las distintas opciones para la función plot
```


```
plot(rulesPruned[1:6], method="graph", control=list(type="items"))
```



Podemos visualizar las reglas como una matriz agrupada. Los antecedentes en las columnas son agrupados usando clustering. En modo interactivo podemos hacer zoom del nodo que queramos estudiar y acceder a las reglas que lo componen para inspeccionarlas.

```
#try: plot(rulesPruned, method="grouped", interactive=TRUE)
plot(rulesPruned[1:6], , method="paracoord", control=list(reorder=TRUE))
```

Parallel coordinates plot for 6 rules



Las podemos guardar en texto plano usando la función write. En este ejemplo las guardamos en un fichero llamado data.csv, usamos como separador “,” y no le ponemos ningún nombre a las columnas

```
write(rulesPruned, file="reglas.csv", sep = ",", col.names=NA)
```

También las podemos guardar en formato PMML. Así podemos volver a leerlas

```
library(pmml)
```

```
## Loading required package: XML
```

```
## Warning: package 'XML' was built under R version 3.3.2
```

```
write.PMML(rulesPruned,file="reglas.pmml")
```

```
## [1] "reglas.pmml"
```

```
reglasPMML = read.PMML("reglas.pmml")
```